Taylor & Francis
Taylor & Francis Group

Journal of the American Statistical Association

ISSN: 0162-1459 (Print) 1537-274X (Online) Journal homepage: www.tandfonline.com/journals/uasa20

PALAR: Estimation of Absolute Abundance Effects in Regression with Relative Abundance Predictors

Yiluan Li, Qiyu Wang, Zekang Feng, Xueqin Wang & Zheng-Zheng Tang

To cite this article: Yiluan Li, Qiyu Wang, Zekang Feng, Xueqin Wang & Zheng-Zheng Tang (02 Mar 2026): PALAR: Estimation of Absolute Abundance Effects in Regression with Relative Abundance Predictors, Journal of the American Statistical Association, DOI: 10.1080/01621459.2025.2596250

To link to this article: https://doi.org/10.1080/01621459.2025.2596250

View supplementary material

Published online: 02 Mar 2026.

Submit your article to this journal

Article views: 351

View related articles

View Crossmark data

Full Terms & Conditions of access and use can be found at
https://www.tandfonline.com/action/journalInformation?journalCode=uasa20

JOURNAL OF THE AMERICAN STATISTICAL ASSOCIATION
2025, VOL. 00, NO. 0, 1–9: Applications and Case Studies
https://doi.org/10.1080/01621459.2025.2596250

Taylor & Francis
Taylor & Francis Group

Check for updates

PALAR: Estimation of Absolute Abundance Effects in Regression with Relative Abundance Predictors

Yiluan Li[a,*], Qiyu Wang[b,*], Zekang Feng[a,*], Xueqin Wang[a], and Zheng-Zheng Tang[c]

[a]School of Management, University of Science and Technology of China, Hefei, Anhui, China; [b]School of Artificial Intelligence and Data Science, University of Science and Technology of China, Hefei, Anhui, China; [c]Department of Biostatistics and Medical Informatics, University of Wisconsin-Madison, Madison, WI

ABSTRACT

High-dimensional compositional data are ubiquitous in omics research. Microbiome sequencing experiments measure the relative abundances (proportions) of microbial features, while the absolute abundances within the ecosystem remain unobserved. Most regression methods with microbial relative abundance predictors rely on log-ratio transformations and typically impose sparsity on the regression coefficients to manage high dimensionality. However, we show that the sparsity assumption often does not hold for these coefficients. To address this issue, we systematically investigate the relationship between the log-ratio regression and the absolute abundance regression. Motivated by the connection, we propose PALAR, a simple and efficient approach for estimating sparse absolute abundance effects using penalized regression with predictors derived from a new compositional transformation. We applied PALAR to four microbiome studies on colorectal cancer, demonstrating its advantages over existing methods in consistently identifying disease-relevant microbial species and improving prediction accuracy and generalizability. Supplementary materials for this article are available online, including a standardized description of the materials available for reproducing the work.

ARTICLE HISTORY
Received April 2025
Accepted November 2025

KEYWORDS
Compositional data; Log-ratio transformation; Microbiome; Penalized regression; Sparse coefficients

1. Introduction

Compositional data, which represent the relative abundances (RA) of components, are prevalent across numerous disciplines, including biology, economics, and ecology. In recent years, high-dimensional compositional datasets have become increasingly common, particularly in omics fields such as microbiome, metabolomics, and proteomics, where relative quantities are often easier to measure than their absolute counterparts (Greenacre, Martínez-Álvaro, and Blasco 2021). Notably, microbiome data generated through high-throughput sequencing consist of sequence read counts assigned to hundreds of microbial features (at the microbe or gene level), which are distorted from their true absolute abundances (AA) within a unit volume of the ecosystem. Since sequencing depth is uncorrelated with the actual microbial load (i.e., the total absolute abundance across all features), the observed counts reflect RA (Gloor et al. 2016, 2017) and are routinely normalized to proportions. Although experimental techniques exist to recover AA microbiome data (Fredricks et al. 2007; Ryu et al. 2013; Vandeputte et al. 2017), they remain costly and their accuracy is subject to debate (McLaren, Willis, and Callahan 2019).

In this article, we focus on regression analysis with RA predictors. In high-dimensional settings, penalized regression is widely used to achieve a balance between predictive accuracy, interpretability, and parsimony. Sparsity is typically induced by applying regularization penalties to the regression coefficients (Tibshirani 1996; Zhu et al. 2022). However, the inherent constraint that RAs must sum to a fixed constant renders traditional penalized regression methods inapplicable (Aitchison 1982). Figure 1 illustrates that AA and RA effects have distinct sparsity profiles. Because a change in the AA of even a single feature alters the RA of all features, sparsity in RA effects is rarely observed (Figure 1(a)). A sparse RA effect profile emerges only when the effects of a small subset of features cancel each other out. For example, in the bottom panels of Figure 1(b), (c), the effects of features A and B balance each other out without involving other features. In this scenario, a corresponding sparse AA effect profile always exists (Figure 1b top panel). Naturally, the underlying AA effect profile is not unique due to the information loss inherent in RA data. Although alternative AA effect profiles exist, they are generally not sparse (Figure 1(c) top panel). The illustration hints at several points that motivate our method: (a) the sparsity assumption on RA effects is generally unrealistic; (b) the sparsity assumption on AA effects is more plausible, interpretable, and of fundamental interest in many applications; and (c) sparse AA effects appear to be approximately identifiable from RA data.

Log-ratio regression is the prevailing approach for regression analysis with RA predictors. In this approach, a log-ratio transformation is applied to the composition to map it from the constrained simplex space to Euclidean space, enabling the use of standard regression techniques. The earliest form of log-ratio regression is based on the additive log-ratio (ALR) transformation (Aitchison and Bacon-shone 1984), with the linear

CONTACT Zheng-Zheng Tang tang@biostat.wisc.edu Department of Biostatistics and Medical Informatics, University of Wisconsin-Madison, Madison, WI.
*The first three authors contributed equally to this work.
Supplementary materials for this article are available online. Please go to www.tandfonline.com/r/JASA.
© 2025 American Statistical Association

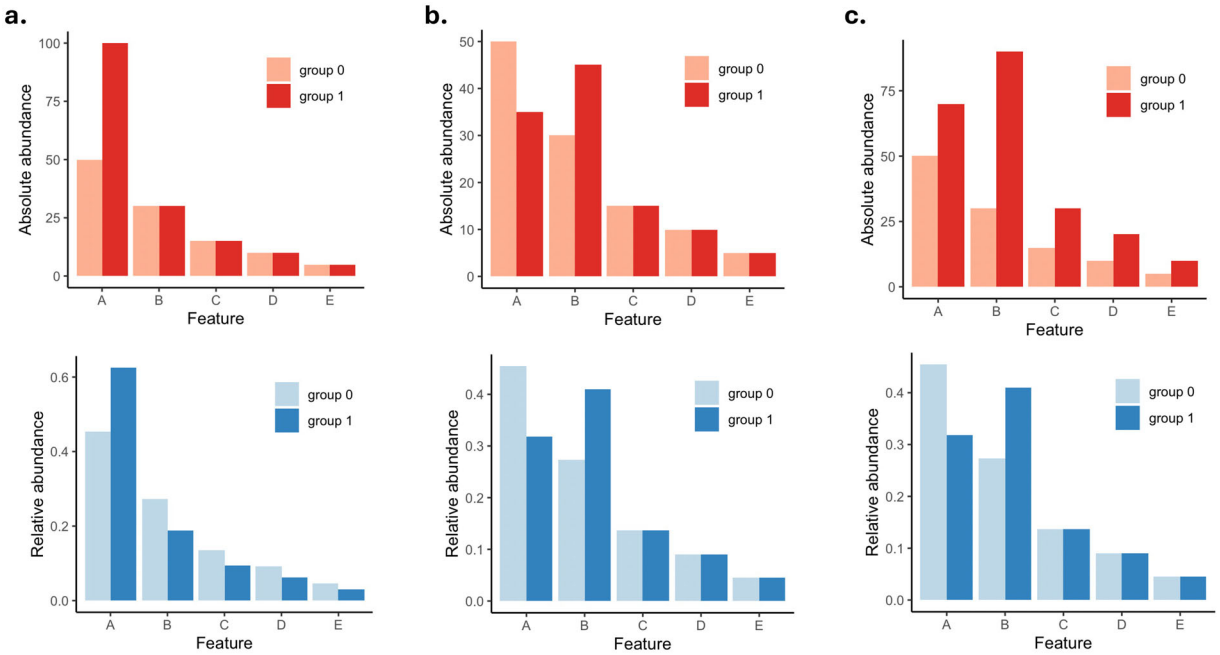

Figure 1. Comparison of effect profiles of absolute abundance (AA, top panel) and relative abundance (RA, bottom panel) predictors in different scenarios **a**–**c**. In **a**, only Feature A has differential AA between comparison groups; in this scenario, the AA effect profile is sparse, but the RA effect profile is not. The RA effect profiles in **b** and **c** are identical and exhibit a sparse pattern because the changes in Features A and B cancel each other out. This may arise from sparse AA effects (**b**, top) or non-sparse AA effects (**c**, top).

predictor: $\beta_1^* \log(Z_1/Z_p) + \cdots + \beta_{p-1}^* \log(Z_{p-1}/Z_p)$, where $(Z_1, \ldots, Z_p)$ is the vector of RAs for p features, and the last feature serves as the reference. To avoid the arbitrariness of selecting a reference component, the log-contrast model (LCM) offers a symmetric alternative: $\beta_1' \log(Z_1) + \cdots + \beta_p' \log(Z_p)$, subject to the zero-sum constraint $\sum_{j=1}^p \beta_j' = 0$. Lin et al. (2014) introduced a variable selection approach for LCM by imposing an ℓ_1 penalty on the coefficients β_j''s. The centered log-ratio transformation is also commonly used, and studies have shown that it is equivalent to the LCM when the zero-sum constraint is imposed on the coefficients (Wang and Zhao 2017; Randolph et al. 2018). Another line of work considers all pairwise log-ratios (Bates and Tibshirani 2019; Greenacre 2019; Gordon-Rodriguez, Quinn, and Cunningham 2022), but these approaches often suffer from computational inefficiencies and identifiability issues.

In this article, we systematically investigate the connection between log-ratio regression and the underlying AA regression, based on the fact that the feature space of any log-ratio-transformed composition lies in a $(p-1)$-dimensional subspace formed by linear combinations of the log-transformed AA variables, representing a one-dimensional reduction from the original space. We elucidate the relationship between AA effects in AA regression and RA effects in log-ratio regression, demonstrating that RA effects are generally non-sparse, even when the underlying AA effects are sparse. We also characterize the information loss caused by this one-dimensional reduction and examine its impact on model compatibility and prediction accuracy. Motivated by the connection between AA and RA effects, we introduce a new compositional transformation and develop PALAR (Penalized AA effect estimation based on a transformed Log-rAtio Regression), a penalized regression framework that uses predictors generated from this new transformation. Under PALAR, we establish the identifiability of sparse AA effects, derive estimation error bounds, and prove variable selection consistency. We demonstrate the advantages of PALAR over existing methods across a wide range of simulation scenarios and apply it to our motivating microbiome datasets of colorectal cancer (as detailed in the next section). Compared to existing methods, PALAR more consistently estimates microbial associations with colorectal cancer across studies, reliably identifies meaningful microbial signatures, and improves colorectal cancer classification accuracy and generalizability.

The remainder of this article is structured as follows. Section 2 introduces motivating datasets and their related scientific questions. Section 3 establishes the connection between log-ratio regression and AA regression, derives the relationship between their coefficients, and characterizes the information loss inherent in log-ratio regression. Section 4 presents the PALAR methodology and establishes theoretical results. Simulation studies and applications to the motivating datasets are presented in Sections 5 and 6, respectively. Section 7 provides a concluding discussion. Proofs and additional results are provided in the supplementary materials.

2. Motivating Datasets and Scientific Questions

Colorectal cancer (CRC) remains one of the leading causes of cancer-related morbidity and mortality worldwide (Sung et al. 2021). In recent years, the study of the gut microbiome has emerged as a promising frontier in CRC research and medicine. The gut microbiome plays a critical role in modulating the immune system, metabolizing dietary compounds, and influencing inflammation—all of which are key factors

in colorectal carcinogenesis (Brennan and Garrett 2016). Our motivating datasets are from four microbiome studies of CRC involving subjects from different countries (Zeller et al. 2014; Vogtmann et al. 2016; Yu et al. 2017; Wirbel et al. 2019). Studies 1–4 had sample sizes of 127 (73 cases, 54 controls), 120 (60 cases, 60 controls), 113 (52 cases, 61 controls), and 103 (51 cases, 52 controls). All microbiome data were generated using metagenomic sequencing. Raw sequencing data were reprocessed using a unified bioinformatics pipeline for taxonomic profiling (Wirbel et al. 2019). The reprocessed data were obtained from *https://github.com/zellerlab/crc_meta*. After filtering out microbial species with less than 20% prevalence across all samples, 401 species were retained for analysis. To handle zero counts, we added a pseudocount of 0.5 to the count matrix.

Our goal is to accurately and reliably identify microbial species whose underlying AA are associated with CRC. The existing penalized regression analyses with RA predictors and ANCOM-BC2 differential abundance test (Lin and Peddada 2024) identify numerous species, most of which exhibit negative effects (Figure S10). Biologically, however, one expects more species to be enriched than depleted in CRC, since tumor-driven alterations create "boom niches" for a handful of organisms rather than causing broad, uniform suppression (Castellarin et al. 2012; Rubinstein et al. 2013). Indeed, meta-analyses consistently report dozens of species enriched in CRC cohorts but only modest depletions of butyrate-producing commensals and other beneficial microbes (Thomas et al. 2019). This pronounced skew toward negative associations likely reflects compositional artifacts rather than true shifts in AA. To address this, we introduce a method that directly estimates sparse AA effects and selects features with genuine nonzero associations. Applying our method yields more positive than negative associations (Figure S10), in line with established CRC–microbiome signatures (details in Section 6).

3. Bridging Log-Ratio Regression to AA Regression

Suppose the data consist of n independent samples. For sample $i = 1, \ldots, n$, let Y_i denote the continuous outcome variable, $\boldsymbol{X}_i = (X_{i1}, \ldots, X_{ip})$ denote the p-dimensional vector of log-transformed AA variables, $\boldsymbol{X}_i^* = (X_{i1}^*, \ldots, X_{ip-1}^*)$ denote the $(p-1)$-dimensional vector of log-ratio-transformed RA variables. Throughout, all covariates are centered across samples. The AA and log-ratio linear regression models are expressed as

$$\text{AA linear regression: } Y_i = \alpha + \boldsymbol{X}_i\boldsymbol{\beta} + \epsilon_i, \tag{1}$$

$$\text{Log-ratio linear regression: } Y_i = \alpha^* + \boldsymbol{X}_i^*\boldsymbol{\beta}^* + \epsilon_i^*, \tag{2}$$

where $\boldsymbol{\beta}$ and $\boldsymbol{\beta}^*$ are the regression coefficients (referred to as AA and RA effects), α and α^* are the intercept coefficients, and ϵ_i and ϵ_i^* are zero-mean normally distributed noise.

For any type of log-ratio transformation, $\boldsymbol{X}_i^*$ can be expressed as a linear projection of $\boldsymbol{X}_i$ onto a $(p-1)$-dimensional subspace. That is, there exists a $p \times (p-1)$ transformation matrix $\boldsymbol{R}$ such that $\boldsymbol{X}_i^* = \boldsymbol{X}_i\boldsymbol{R}$. Based on this connection, we examine the relationship between the two models.

We view $\{\boldsymbol{X}_i\}_{i=1}^n$ as i.i.d. draws of a zero-mean random vector with covariance matrix $\boldsymbol{\Sigma}_X = E(\boldsymbol{X}_i^\top \boldsymbol{X}_i)$ (referred to as the AA covariance matrix).

3.1. Relationship between AA and RA Effects

Let $\boldsymbol{X} = (\boldsymbol{X}_1^\top, \ldots, \boldsymbol{X}_n^\top)^\top$ and $\boldsymbol{X}^* = \boldsymbol{X}\boldsymbol{R}$. The linear predictors in the two models have a relationship

$$\boldsymbol{X}\boldsymbol{\beta} = \boldsymbol{X}^*\boldsymbol{\beta}^* + \boldsymbol{\xi}, \tag{3}$$

where $\boldsymbol{\xi}$ lies in the orthogonal complement of the column space of $\boldsymbol{X}^*$.

Therefore, the RA effect $\boldsymbol{\beta}^*$ is given by the population least-squares projection coefficient

$$\boldsymbol{\beta}^* = \left\{E(\boldsymbol{X}_i^{*\top}\boldsymbol{X}_i^*)\right\}^{-1} E(\boldsymbol{X}_i^{*\top}\boldsymbol{X}_i)\boldsymbol{\beta} = (\boldsymbol{R}^\top\boldsymbol{\Sigma}_X\boldsymbol{R})^{-1}\boldsymbol{R}^\top\boldsymbol{\Sigma}_X\boldsymbol{\beta}. \tag{4}$$

Clearly, $\boldsymbol{\beta}^*$ generally lacks sparsity, even when $\boldsymbol{\beta}$ is sparse. Next, we take the ALR transformation as a concrete example to further illustrate the relationship between $\boldsymbol{\beta}$ and $\boldsymbol{\beta}^*$.

Proposition 3.1. In ALR regression, suppose the pth feature is used as the reference. Assuming this reference is uncorrelated with the other features in the AA space, $\boldsymbol{\beta}^*$ takes the form

$$\boldsymbol{\beta}^* = \boldsymbol{\beta}_{\backslash p} - C\boldsymbol{\Sigma}_{X_{\backslash p}}^{-1}\mathbf{1}_{p-1},$$

$$C = \frac{\mathbf{1}_p^\top\boldsymbol{\beta}}{\mathbf{1}_p^\top\boldsymbol{\Sigma}_X^{-1}\mathbf{1}_p},$$

where $\boldsymbol{\beta}_{\backslash p}$ denotes the subvector of $\boldsymbol{\beta}$ excluding its pth element, $\boldsymbol{\Sigma}_{X_{\backslash p}}$ denotes the submatrix of $\boldsymbol{\Sigma}_X$ removing the pth row and column, and $\mathbf{1}_p$ denotes the p-dimensional vector of ones.

The proof of Proposition 3.1 is provided in the supplementary materials. Note that the scalar C is invariant to the choice of reference feature and increases with the sum of the AA effects $\mathbf{1}_p^\top\boldsymbol{\beta}$, which tends to be large when most active features share the same direction of effect. The term $C\boldsymbol{\Sigma}_{X_{\backslash p}}^{-1}\mathbf{1}_{p-1}$ depends on the choice of the reference and induces a systematic difference between $\boldsymbol{\beta}$ and $\boldsymbol{\beta}^*$.

3.2. Information Loss

The term $\boldsymbol{\xi}$ in (3) represents the inherent information loss in log-ratio regression relative to the AA regression. The following proposition outlines several properties of $\boldsymbol{\xi}$.

Proposition 3.2. The information loss $\boldsymbol{\xi}$ satisfies the following properties:

(a) $\boldsymbol{\xi}$ is uncorrelated with the linear predictor in log-ratio regression and invariant to the type of log-ratio transformation.
(b) If $\boldsymbol{X}$ follows a multivariate normal distribution, then $\boldsymbol{\xi}$ is also normally distributed with mean $\mathbf{0}$ and variance $\boldsymbol{\beta}^\top\boldsymbol{\Sigma}_X\boldsymbol{Q}_R\boldsymbol{\beta}$, where $\boldsymbol{Q}_R = \boldsymbol{I}_p - \boldsymbol{R}(\boldsymbol{R}^\top\boldsymbol{\Sigma}_X\boldsymbol{R})^{-1}\boldsymbol{R}^\top\boldsymbol{\Sigma}_X$.
(c) $\boldsymbol{\xi} \approx (\mathbf{1}_p^\top\boldsymbol{\beta})(\boldsymbol{X}\mathbf{1}_p^\top/p)$

The proof of Proposition 3.2 is provided in the supplementary materials. If $\boldsymbol{X}$ is multivariate normally distributed, the log-ratio linear regression (2) is compatible with the AA linear regression (1), in the sense that the information loss is

captured in the residual ϵ_i^*. From part (c) of the proposition, we know that the information loss increases with the sum of the AA effects. When the sum of the AA effects is zero, the log-ratio regression incurs no information loss relative to the AA regression. Figure S1 displays simulation results comparing the prediction errors of the two regressions. As expected, the mean squared error (MSE) of the log-ratio regression increases quadratically as the imbalance between positive and negative AA effects grows. It is important to emphasize that it is not possible to recover this information loss and close the prediction performance gap between AA and RA predictors. Our focus is on efficiently estimating the sparse AA effects, thereby maximizing prediction accuracy within the RA-based log-ratio regression framework, as reflected by achieving the performance indicated by the dashed line in Figure S1.

Beyond linear regression, information loss renders the regular log-ratio model incompatible with the AA model under generalized linear models (GLM). A log-ratio model that is compatible with the AA regression should take the form

$$g[E(Y_i \mid \boldsymbol{X}_i^*, \xi_i)] = \alpha^* + \boldsymbol{X}_i^* \boldsymbol{\beta}^* + \xi_i, \tag{5}$$

where g is the link function. According to Proposition 3.2, ξ_i represents the information loss, which follows a zero-mean normal distribution with variance $\sigma_\xi^2(\boldsymbol{\beta}) = \boldsymbol{\beta}^\top \boldsymbol{\Sigma}_X \boldsymbol{Q}_R \boldsymbol{\beta}$. Since ξ_i is unobserved, inference should be based on the marginal density of Y_i, defined as $f(Y_i \mid \boldsymbol{X}_i^*) = \int_{-\infty}^{\infty} f(Y_i \mid \boldsymbol{X}_i^*, \xi_i)\phi(\xi_i)\, d\xi_i$, where ϕ is the density function of a zero-mean normal distribution with variance $\sigma_\xi^2(\boldsymbol{\beta})$. In general, the marginal density does not have a closed-form expression and must be approximated numerically. Iterative algorithms commonly used for parameter estimation in mixed-effects models can be adapted to fit this model.

4. Methodology

4.1. PALAR Linear Regression

As shown in the previous section, the RA effect vector $\boldsymbol{\beta}^*$ is generally not sparse, even when the underlying AA effect vector $\boldsymbol{\beta}$ is sparse. This motivates us to estimate the sparse $\boldsymbol{\beta}$ directly, rather than imposing sparsity on $\boldsymbol{\beta}^*$. To this end, we introduce a new compositional transformation

$$\boldsymbol{X}^{**} = \boldsymbol{X}^*(\boldsymbol{R}^\top \boldsymbol{\Sigma}_X \boldsymbol{R})^{-1} \boldsymbol{R}^\top \boldsymbol{\Sigma}_X. \tag{6}$$

Based on (4), the linear predictor of the log-ratio regression can be written as $\boldsymbol{X}^{**}\boldsymbol{\beta}$. The transformed variables $\boldsymbol{X}^{**}$ have several nice properties and interpretations:

- $\boldsymbol{X}^{**}$ has the highest overall correlation with $\boldsymbol{X}$, defined as the trace of the covariance $\mathrm{cov}(\boldsymbol{X}, \boldsymbol{X}^{**})$, among all linear transformations of $\boldsymbol{X}^*$.
- $\boldsymbol{X}^{**}$ is invariant to the choice of log-ratio transformation. That is, for any two log-ratio transformation matrices $\boldsymbol{R}_1$ and $\boldsymbol{R}_2$, we obtain the same result: $\boldsymbol{X}^{**} = \boldsymbol{X}\boldsymbol{R}_1(\boldsymbol{R}_1^\top \boldsymbol{\Sigma}_X \boldsymbol{R}_1)^{-1}\boldsymbol{R}_1^\top \boldsymbol{\Sigma}_X = \boldsymbol{X}\boldsymbol{R}_2(\boldsymbol{R}_2^\top \boldsymbol{\Sigma}_X \boldsymbol{R}_2)^{-1}\boldsymbol{R}_2^\top \boldsymbol{\Sigma}_X$.
- $\boldsymbol{X}^{**}$ recovers the AA effects in the linear predictor. Specifically, for any AA effect vector $\boldsymbol{\beta}$, the linear predictors based on $\boldsymbol{X}$ and $\boldsymbol{X}^{**}$ satisfy: $\boldsymbol{X}\boldsymbol{\beta} = \boldsymbol{X}^{**}\boldsymbol{\beta} + \boldsymbol{\xi}$. In contrast, the predictor effect in the log-ratio regression is different from the AA effects: $\boldsymbol{X}\boldsymbol{\beta} = \boldsymbol{X}^*\boldsymbol{\beta}^* + \boldsymbol{\xi}$ and $\boldsymbol{\beta}^* \neq \boldsymbol{\beta}$.

The following theorem establishes the condition under which sparse $\boldsymbol{\beta}$ is identifiable using Lasso regression with predictors $\boldsymbol{X}^{**}$.

Theorem 4.1. The S-sparse AA effect vector $\boldsymbol{\beta}$ is identifiable if the Restricted Isometry Property (RIP) constant of order $2s$ satisfies $\delta_{2s}(\boldsymbol{X}^{**}) < 1/3$.

The condition mirrors that of Proposition 7.11 of Wainwright (2019). Although $\boldsymbol{X}^{**}$ is a projection of $\boldsymbol{X}$ onto a $(p-1)$-dimensional space, its compatibility conditions (e.g., RIP) are preserved (justification provided in the supplementary materials A.3). As a result, $\boldsymbol{X}^{**}$ remains effective for ensuring the identifiability of $\boldsymbol{\beta}$.

The $\boldsymbol{X}^{**}$ involves the unknown AA covariance matrix $\boldsymbol{\Sigma}_X$. Several methods have been proposed in the literature to estimate $\boldsymbol{\Sigma}_X$ in the context of high-dimensional compositional data. In our numerical studies, we adopt the SparCC method (Friedman and Alm 2012). Let $\widetilde{\boldsymbol{X}}^{**}$ denote the surrogate of $\boldsymbol{X}^{**}$ obtained by plugging in an estimate of $\boldsymbol{\Sigma}_X$. The PALAR ℓ_1-penalized estimate of $\boldsymbol{\beta}$ is defined as the solution to the optimization problem

$$(\widehat{\alpha}^*, \widehat{\boldsymbol{\beta}}) = \arg\min \left\{ \frac{1}{2n} \sum_{i=1}^{n} (Y_i - \alpha^* - \widetilde{\boldsymbol{X}}_i^{**}\boldsymbol{\beta})^2 + \lambda_n \|\boldsymbol{\beta}\|_1 \right\}. \tag{7}$$

This problem can be efficiently solved using standard Lasso solvers and the tuning parameter λ_n can be selected via cross-validation.

We next present theoretical results on the estimation error bounds and variable selection consistency of the PALAR estimator $\widehat{\boldsymbol{\beta}}$. These results are adapted from standard Lasso theory (Wainwright 2019), with modifications to account for the approximation error introduced by using the surrogate $\widetilde{\boldsymbol{X}}^{**}$ in place of the true $\boldsymbol{X}^{**}$. Proofs are provided in the supplementary materials. Let $\boldsymbol{Y}$, $\boldsymbol{X}^{**}$, $\widetilde{\boldsymbol{X}}^{**}$, and $\boldsymbol{\epsilon}^*$ denote the stacked vectors or matrices of Y_i, $\boldsymbol{X}_i^{**}$, $\widetilde{\boldsymbol{X}}_i^{**}$, and ϵ_i^*, respectively. Let $\boldsymbol{\beta}^0$ denote the true AA effect vector, S be its support set of cardinality s, $\widetilde{\boldsymbol{w}} = (\boldsymbol{X}^{**} - \widetilde{\boldsymbol{X}}^{**})\boldsymbol{\beta}^0$, and $\widetilde{\boldsymbol{X}}_S^{**}$ denote the submatrix of $\widetilde{\boldsymbol{X}}^{**}$ consisting of the columns corresponding to the indices in the set S. For any vector $\boldsymbol{v}$, let $\|\boldsymbol{v}\|_r$ be the vector ℓ_r-norm. For simplicity, we omit the intercept term in the results below.

Theorem 4.2. Consider the S-sparse $\boldsymbol{\beta}^0$ in the linear regression model $\boldsymbol{Y} = \boldsymbol{X}^{**}\boldsymbol{\beta} + \boldsymbol{\epsilon}^*$, if the surrogate design matrix $\widetilde{\boldsymbol{X}}^{**}$ satisfies the restricted eigenvalue (RE) condition with parameters $(\kappa, 3)$ and the regularization parameter lower bounded as $\lambda_n \geq 2\|\boldsymbol{X}^{**\top}\boldsymbol{\epsilon}^*/n\|_\infty + 2\|\widetilde{\boldsymbol{X}}^{**\top}\widetilde{\boldsymbol{w}}/n\|_\infty$, the PALAR estimator $\widehat{\boldsymbol{\beta}}$ obtained from (7) satisfies the bound

$$\|\widehat{\boldsymbol{\beta}} - \boldsymbol{\beta}^0\|_2 \leq \frac{3}{\kappa}\sqrt{s}\lambda_n.$$

This theorem ensures that the ℓ_2 estimation error of the PALAR Lasso is small. The lower bound on λ_n differs from standard Lasso theory due to the need to account for the difference between $\widetilde{\boldsymbol{X}}^{**}$ and $\boldsymbol{X}^{**}$. As $\widetilde{\boldsymbol{X}}^{**}$ better approximates $\boldsymbol{X}^{**}$, the term $\|\widetilde{\boldsymbol{X}}^{**\top}\widetilde{\boldsymbol{w}}/n\|_\infty$ becomes smaller, thereby reducing the overall estimation error. Corollary of this theorem, given an assumption

on $\|\boldsymbol{\Sigma}_X - \widetilde{\boldsymbol{\Sigma}}_X\|_\infty$, was presented in the supplementary materials A.5.

Theorem 4.3. Consider S-sparse $\boldsymbol{\beta}^0$ in the linear regression model $\boldsymbol{Y} = \boldsymbol{X}^{**}\boldsymbol{\beta} + \boldsymbol{\epsilon}^*$. Suppose the surrogate design matrix $\widetilde{\boldsymbol{X}}^{**}$ satisfies the conditions

(a) The smallest eigenvalue of the Gram matrix satisfies $\gamma_{\min}\left(\frac{\widetilde{\boldsymbol{X}}_S^{**\top}\widetilde{\boldsymbol{X}}_S^{**}}{n}\right) \geq C_{\min} > 0$.
(b) There exists $a \in [0,1]$ such that $\max_{j\in S^c} \|(\widetilde{\boldsymbol{X}}_S^{**\top}\widetilde{\boldsymbol{X}}_S^{**})^{-1}\widetilde{\boldsymbol{X}}_S^{**\top}\widetilde{\boldsymbol{X}}_j^{**}\|_1 \leq a$,

Then for any regularization parameter $\lambda_n \geq \frac{2}{1-a}\left\|\widetilde{\boldsymbol{X}}_{S^c}^{**\top}\Pi_{S^\perp}(\widetilde{\boldsymbol{X}}^{**})\frac{\boldsymbol{\epsilon}^*+\widetilde{\boldsymbol{w}}}{n}\right\|_\infty$, where $\Pi_{S^\perp}(\widetilde{\boldsymbol{X}}_S^{**}) = \boldsymbol{I}_n - \widetilde{\boldsymbol{X}}_S^{**}\left(\widetilde{\boldsymbol{X}}_S^{**\top}\widetilde{\boldsymbol{X}}_S^{**}\right)^{-1}\widetilde{\boldsymbol{X}}_S^{**\top}$, the PALAR Lasso (7) has the following properties:

(1) There is a unique optimal solution $\widehat{\boldsymbol{\beta}}$.
(2) This solution has its support set $\widehat{S}$ contained within the true support set S.
(3) $\|\widehat{\boldsymbol{\beta}}_S - \boldsymbol{\beta}_S^0\|_\infty \leq \left\|\left(\frac{\widetilde{\boldsymbol{X}}_S^{**\top}\widetilde{\boldsymbol{X}}_S^{**}}{n}\right)^{-1}\widetilde{\boldsymbol{X}}_S^{**\top}\frac{\boldsymbol{\epsilon}^*+\widetilde{\boldsymbol{w}}}{n}\right\|_\infty + \left|\left|\left|\left(\frac{\widetilde{\boldsymbol{X}}_S^{**\top}\widetilde{\boldsymbol{X}}_S^{**}}{n}\right)^{-1}\right|\right|\right|_\infty \lambda_n$, where $|||\boldsymbol{X}|||_\infty$ is the matrix ℓ_∞ norm.
(4) No false exclusion if $\min_{i\in S}|\beta_i^0| > B(\lambda_n; \widetilde{\boldsymbol{X}}^{**})$, where $B(\lambda_n; \widetilde{\boldsymbol{X}}^{**})$ is the ℓ_∞ upper bound in (3).

Property (2) states that PALAR Lasso does not falsely include variables outside the true support of $\boldsymbol{\beta}^0$. Property (4) ensures it does not falsely exclude variables within the support. These properties establish that PALAR Lasso achieves variable selection consistency. Corollary of this theorem, given an assumption on $\|\boldsymbol{\Sigma}_X - \widetilde{\boldsymbol{\Sigma}}_X\|_\infty$, was presented in the supplementary materials A.7.

4.2. PALAR Generalized Linear Regression

From Section 3.2, we know that the regular log-ratio GLM and the AA GLM are not compatible. Although it would be desirable to work with the compatible model given in (5), the associated computational burden is nontrivial and will be addressed in future work.

In this article, we proceed by ignoring the information loss in the model. Specifically, in logistic regression with a binary outcome Y_i, the PALAR ℓ_1-penalized estimate of $\boldsymbol{\beta}$ is defined as the solution to the optimization problem:

$$(\widehat{\alpha}^*, \widehat{\boldsymbol{\beta}}) = \arg\min\left\{-\frac{1}{n}\sum_{i=1}^n \ell_i(\alpha^*, \boldsymbol{\beta}) + \lambda_n\|\boldsymbol{\beta}\|_1\right\}, \tag{8}$$

$$\text{where } \ell_i(\alpha^*, \boldsymbol{\beta}) = Y_i \log\left(\frac{e^{\alpha^*+\widetilde{\boldsymbol{X}}_i^{**}\boldsymbol{\beta}}}{1+e^{\alpha^*+\widetilde{\boldsymbol{X}}_i^{**}\boldsymbol{\beta}}}\right) + (1-Y_i)\log\left(\frac{1}{1+e^{\alpha^*+\widetilde{\boldsymbol{X}}_i^{**}\boldsymbol{\beta}}}\right).$$

This optimization problem can be efficiently solved using standard Lasso solvers and the tuning parameter λ_n can be selected via cross-validation.

Although the information loss is not explicitly addressed, simulation studies and real data applications show that PALAR substantially outperforms regular log-ratio logistic regression in terms of estimation accuracy, prediction performance, and variable selection.

5. Simulation Studies

We conducted simulation studies to evaluate the numerical performance of PALAR using the estimated AA covariance and the true AA covariance (PALAR-TC). We compared these methods with ALR regression, LCM, and the regression model using log-proportion predictors without the zero-sum constraint on the regression coefficients (LP). In ALR regression, the reference feature was randomly selected from the p features. LCM was applied only in the continuous outcome setting due to the availability of R code from Lin et al. (2014). All other methods were implemented using the Lasso solver from the glmnet R package. All methods used 10-fold cross-validation to select the tuning parameter.

We simulated log(AA) data $\boldsymbol{X} = \{x_{ij}\}$ from a multivariate normal distribution $\mathcal{N}_p(\boldsymbol{\mu}_X, \boldsymbol{\Sigma}_X)$, and then obtained the proportions $\boldsymbol{P} = \{p_{ij}\}$ using the softmax transformation $p_{ij} = \exp(x_{ij})/\sum_{k=1}^p \exp(x_{ik})$. The parameters $\boldsymbol{\mu}_X$ and $\boldsymbol{\Sigma}_X$ were specified based on one of the motivating datasets as the template. Specifically, we first extracted the count matrix for p randomly selected features from the template. Then, we applied SparCC to the count matrix to obtain the AA covariance estimate and defined it as $\boldsymbol{\Sigma}_X$. The sample mean of the log-transformed count matrix was used to define $\boldsymbol{\mu}_X$.

Using the simulated $\boldsymbol{X}$, we generated continuous outcomes according to the AA linear regression model and binary outcomes according to the AA logistic regression model. We randomly selected 10% of the features as active features with nonzero AA effects. To evaluate the impact of effect direction imbalance, we considered three scenarios: $d = 0.5$ represents an equal number of positive and negative effects; $d = 0.75$ represents that 75% of the active features have negative effects and 25% have positive effects; $d = 1$ corresponds to all active features having negative effects. The magnitudes of the nonzero effects were sampled from $|N(0,1)|$.

We considered sample sizes $n = 200, 400, 800$ and feature dimensions $p = 200, 400$, and repeated 100 simulations for each setting. Performance was evaluated using four metrics. Estimation error was measured by the ℓ_1 loss $\|\widehat{\boldsymbol{\beta}} - \boldsymbol{\beta}^0\|_1$, where $\widehat{\boldsymbol{\beta}}$ denotes the estimated effects. In ALR regression, the coefficient for the reference feature cannot be estimated and is therefore fixed at zero. Prediction performance was assessed using MSE for continuous outcomes and the area under the receiver operating characteristic curve (AUC) for binary outcomes. Both metrics were evaluated on an independent test set of size n. To construct the PALAR predictors $\boldsymbol{X}^{**}$ in the test set, we used the AA covariance matrix $\boldsymbol{\Sigma}_X$ estimated from the training data. We also evaluated variable selection performance using the false positive rate (FPR) and false negative rate (FNR), where positives and negatives refer to features with nonzero and zero AA effects, respectively.

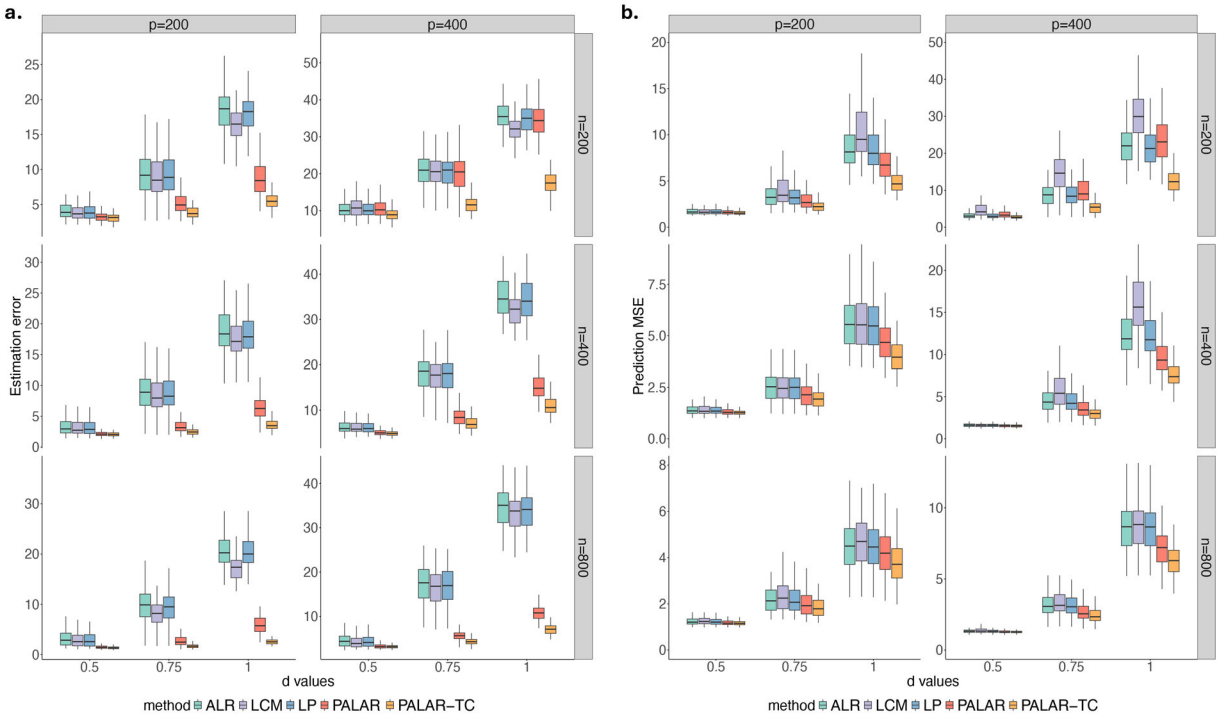

Figure 2. Comparison of different methods based on estimation error (**a**) and prediction error (**b**) in simulations with the continuous outcome under varying *d* values. The *d* value is the proportion of active features with negative effects.

Figure 2 presents the results on estimation and prediction performance, and Figure S2 shows the variable selection performance. As d increases, the effect directions become more imbalanced, and the sum of effect sizes deviates further from zero. In the balanced scenario ($d = 0.5$), the performance of all methods is relatively similar. However, as the imbalance in effect directions increases, the advantages of PALAR-TC and PALAR become more pronounced. PALAR-TC consistently outperforms PALAR, as it uses the true covariance matrix; this performance gap narrows with increasing sample size, reflecting improved estimation of the AA covariance matrix. ALR and LP regressions perform similarly across all settings. Compared to ALR and LP regressions, LCM yields more accurate estimation but worse prediction performance, with lower FPR but higher FNR. Figure S3 provides an illustrative example to further highlight the variable selection performance of different methods under varying values of d. The estimated regression coefficients vary depending on the specific transformation applied to the compositional predictors. As d increases, ALR and LP regressions and LCM tend to select a large number of null features. In contrast, PALAR achieves more accurate selections and demonstrates robustness to changes in d.

The results of the logistic regression simulations are shown in Figures S4 and S5. PALAR and PALAR-TC continue to outperform the other methods, with their advantages becoming more pronounced as d increases. In addition, we conducted a separate set of simulation studies that incorporated noise in the RA data to reflect measurement errors commonly observed in real microbiome datasets. Specifically, we simulated sampling count data with varying proportions of zeros, examined the impact of different pseudocount choices when converting count data to noisy RA for analysis, and assessed sensitivity to the assumption of a multivariate normal distribution for log(AA). We also compared variable selection with ANCOM-BC2. PALAR remained competitive across these simulations (Figures S6–S8).

6. Data analysis

We applied PALAR to the four gut microbiome studies to identify microbial signatures associated with CRC. As in the simulation studies, we also applied ALR and LP regressions and ANCOM-BC2. For ALR, the reference was randomly selected from the 50 species with relatively high abundance and low variance, measured by the coefficient of variation. For ANCOM-BC2, features were selected at a false discovery rate threshold of 0.05.

We first analyzed each study individually and evaluated the concordance of effect estimates across studies. Specifically, for each pair of studies s and s', we computed the sum of squared deviation (SSD) of effect estimates as $\sum_{j=1}^{p}\left(\widehat{\beta}_j^{(s)} - \widehat{\beta}_j^{(s')}\right)^2$, where $\widehat{\beta}_j^{(s)}$ denotes the effect estimate for feature j in study s. To ensure robustness, we repeated this evaluation on 100 bootstrap samples of the four studies and reported the average SSD across replicates (Figure 3). PALAR achieved the highest overall consistency of effect estimates, while LP showed the lowest. This result suggests that PALAR is better able to capture stable microbial associations that are less influenced by inter-study variability and technical batch effects commonly present in microbiome data.

Next, using the pooled dataset from all studies, we identified microbial features associated with CRC and assessed their classification performance. To ensure stable results, we generated 100 bootstrap samples of the pooled dataset. Within

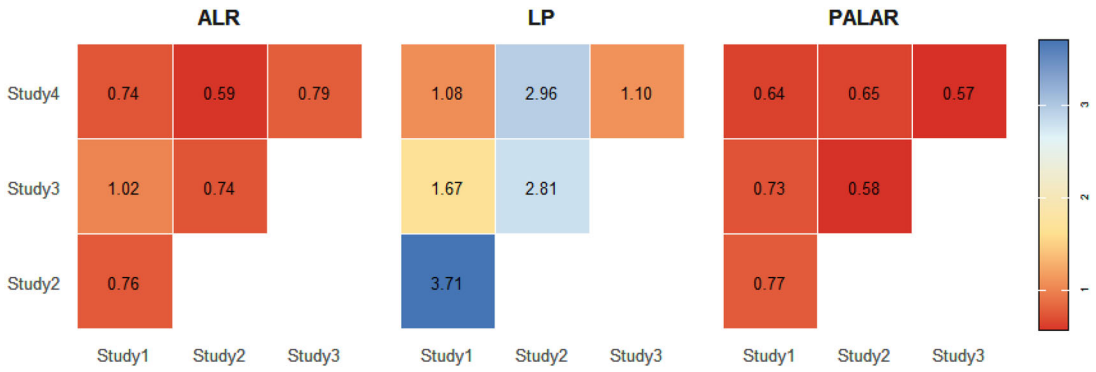

Figure 3. Comparison of methods in effect estimate concordance across studies. The heatmap displays the sum of squared deviation of effect estimates between study pairs, with lower values indicating higher concordance.

each bootstrap replicate, we randomly split the data into a 90% training set for feature selection and a 10% testing set for evaluating prediction model. Species identified by PALAR in at least 90% of the replicates are shown in Table 1. Those selected by ALR and LP regressions and ANCOM-BC2 are shown in Tables S1, S2, and S3, respectively. The results from ALR regression and ANCOM-BC2 are more sensitive to the choice of pseudocount compared to PALAR (Figure S9). PALAR identified several well-established CRC-associated species, including *Fusobacterium nucleatum*, *Peptostreptococcus stomatis*, *Parvimonas micra*, *Gemella morbillorum*, and *Alistipes onderdonkii*. These microbes promote CRC through diverse, well-characterized mechanisms. *Fusobacterium nucleatum* adheres to colonic epithelial cells via its FadA adhesin, activating β-catenin signaling through E-cadherin binding to drive inflammation and proliferation (Rubinstein et al. 2013). *Peptostreptococcus stomatis* promotes tumorigenesis by using its surface FBA to bind integrin $\alpha 6/\beta 4$ on CRC cells, triggering ERBB2–MAPK signaling that enhances epithelial growth and disrupts barrier integrity (Huang et al. 2024). *Parvimonas micra* induces colonocyte proliferation and shifts the mucosal immune profile toward Th17 dominance, amplifying Wnt signaling and inflammation (Zhao et al. 2022). *Gemella morbillorum* fosters a tumor-supportive niche by driving IL-17 and IL-23 production, reinforcing Th17-mediated inflammation (Jobin 2017). *Alistipes onderdonkii* activates the IL-6/STAT3 axis in murine models, promoting tumor cell survival and epithelial proliferation (Parker et al. 2020). Notably, none of the compared methods detected *Alistipes onderdonkii*. Moreover, species selected by PALAR yielded the highest prediction accuracy (Table 2).

We also evaluated the reliability of our results by repeating feature selection on each leave-one-study-out (LOSO) dataset. PALAR's selections are biologically more plausible—consistently identifying more CRC-enriched species than depleted ones—whereas ALR and LP regressions and ANCOM-BC2 exhibit the opposite pattern (Figure S10). In addition, 86% of the species chosen by PALAR (Table 1) were recovered in at least three of the four LOSO analyses, underscoring the method's robustness.

We further performed the LOSO analysis to assess prediction generalizability. The data from one study was set aside as an external testing set, while the data from the remaining three studies was pooled as a training set, on which we performed cross-validation for model tuning. We evaluated the performance of the trained model on the held-out study. We repeated the procedure 100 times with different random splits of the cross-validation folds and reported the results in Table S4. Species selected by PALAR consistently achieved the best CRC

Table 1. CRC-associated gut microbial species identified by PALAR.

Species	Beta
Positive Effects:	
Solobacterium moorei [0531]	0.7484
Peptostreptococcus stomatis [4614]	0.5568
Fusobacterium nucleatum species [0776]	0.5056
unknown Ruminococcaceae [7652]	0.4568
Coprococcus catus [4874]	0.3940
unknown Clostridiales [6348]	0.3770
unknown Anaerotruncus [6835]	0.3667
Gemella morbillorum [4513]	0.3588
Parvimonas micra [1145]	0.3324
Subdoligranulum species [4738]	0.3073
Veillonella parvula [1042]	0.2829
unknown Bacteroidales [5329]	0.2748
unknown Prevotella [5668]	0.2677
Clostridium symbiosum [1475]	0.2385
Acidaminococcus intestini [0391]	0.2033
Alistipes onderdonkii [0775]	0.1952
Ruminococcus species [6664]	0.1893
unknown Clostridiales [6105]	0.1618
unknown Firmicutes [6091]	0.1515
Phascolarctobacterium species [2805]	0.1129
Negative Effects:	
uncultured Eubacterium species [6201]	−0.9617
Blautia producta [4020]	−0.5336
unknown Clostridiales [6107]	−0.4550
unknown Clostridiales [5880]	−0.3247
unknown Clostridiales [6700]	−0.2970
Clostridium species [5681]	−0.2428
unknown Peptostreptococcaceae [7331]	−0.1717
Clostridium species [7788]	−0.1354
unknown Clostridium [5413]	−0.1310

NOTE: Species listed were selected in 90% of bootstrap replicates. Beta is the refitted effect estimate based on the pooled dataset.

Table 2. Prediction performance on bootstrap pooled data.

Method	AUC	Brier
ALR	0.8016 (0.0575)	0.1806 (0.0245)
LP	0.8070 (0.0592)	0.1787 (0.0251)
PALAR	0.8127 (0.0578)	0.1761 (0.0243)

NOTE: Average AUC and Brier scores across 100 bootstrap replicates are reported, with standard deviations shown in parentheses.

prediction performance across all held-out studies, demonstrating strong generalizability.

7. Discussion

Our work is motivated by the analysis of gut microbiome data from CRC studies. A key statistical challenge in such datasets is that microbiome sequencing provides only RA data, while the true biological signals of interest are often encoded in the unobserved AA. Existing regression methods with RA predictors typically rely on log-ratio transformations and impose sparsity assumptions on the coefficients. However, we show that the sparsity assumption often does not hold for these log-ratio regression coefficients, leading to unreliable effect estimation and variable selection. To address this issue, we propose PALAR, a new penalized regression framework that explicitly connects RA-based log-ratio regression to the underlying AA regression. PALAR estimates sparse AA effects using a new compositional transformation that recovers the linear predictor up to an irreducible information loss. We establish theoretical guarantees for identifiability, estimation error, and variable selection consistency.

Extensive simulations demonstrate PALAR's superior estimation, variable selection, and prediction performance, particularly in settings with imbalanced effect directions where existing methods struggle. Analyses of gut microbiome data from CRC studies further demonstrate the advantages of PALAR. The method produced more consistent effect estimates across studies and successfully identified several well-established CRC-associated microbial species that were missed by other methods. The improved concordance of effect estimates and enhanced predictive performance underscore PALAR's potential for uncovering true associations in high-dimensional compositional data.

PALAR offers several avenues for extension and improvement. First, its performance could be enhanced with a more accurate estimation of the AA covariance matrix. Future work may explore more robust and adaptive strategies for estimating this matrix for high-dimensional compositional data. Second, although PALAR has been demonstrated in logistic regression models, the current formulation does not explicitly account for the variation in information loss. Extending PALAR to a generalized linear mixed-effects framework that incorporates this information loss can potentially improve the estimation efficiency. Finally, the current implementation employs the Lasso penalty to induce sparsity. However, PALAR is flexible and agnostic to the choice of regularization. Incorporating phylogenetic tree information into microbiome data analysis has been shown to improve the performance of various statistical methods (Tang et al. 2017; Wang and Zhao 2017; Hong, Chen, and Tang 2023). Our framework can readily accommodate alternative penalties, including structured sparsity approaches such as tree-guided regularization (Chen et al. 2015).

Supplementary Materials

The supplementary materials contain the proofs of theoretical results, additional numerical results, and R code and data.

Acknowledgments

The authors thank the editor, the associate editor, and referees for their constructive feedback, which led to significant improvements in the article.

Disclosure Statement

The authors report there are no competing interests to declare.

Funding

Tang was supported by NIH R01AI192972, NIH R01GM140464, and NSF DMS-2054346.

ORCID

Xueqin Wang https://orcid.org/0000-0001-5205-9950
Zheng-Zheng Tang https://orcid.org/0000-0003-3802-8087

References

Aitchison, J. (1982), "The Statistical Analysis of Compositional Data," *Journal of the Royal Statistical Society*, Series B, 44, 139–177. [1]

Aitchison, J., and Bacon-shone, J. (1984), "Log Contrast Models for Experiments with Mixtures," *Biometrika*, 71, 323–330. [1]

Bates, S., and Tibshirani, R. (2019), "Log-Ratio Lasso: Scalable, Sparse Estimation for Log-Ratio Models," *Biometrics*, 75, 613–624. [2]

Brennan, C. A., and Garrett, W. S. (2016), "Gut Microbiota, Inflammation, and Colorectal Cancer," *Annual Review of Microbiology*, 70, 395–411. [3]

Castellarin, M., Warren, R. L., Freeman, J. D., Dreolini, L., Krzywinski, M., Strauss, J., Barnes, R., Watson, P., Allen-Vercoe, E., Moore, R. A., et al. (2012), "Fusobacterium Nucleatum Infection is Prevalent in Human Colorectal Carcinoma," *Genome Research*, 22, 299–306. [3]

Chen, L., Liu, H., Kocher, J.-P. A., Li, H., and Chen, J. (2015), "glmgraph: An r Package for Variable Selection and Predictive Modeling of Structured Genomic Data," *Bioinformatics*, 31, 3991–3993. [8]

Fredricks, D. N., Fiedler, T. L., Thomas, K. K., Oakley, B. B., and Marrazzo, J. M. (2007), "Targeted PCR for Detection of Vaginal Bacteria Associated with Bacterial Vaginosis," *Journal of Clinical Microbiology*, 45, 3270–3276. [1]

Friedman, J., and Alm, E. J. (2012), "Inferring Correlation Networks from Genomic Survey Data," *PLoS Computational Biology*, 8, e1002687. [4]

Gloor, G. B., Macklaim, J. M., Pawlowsky-Glahn, V., and Egozcue, J. J. (2017), "Microbiome Datasets are Compositional: And This Is Not Optional," *Frontiers in Microbiology*, 8, 2224. [1]

Gloor, G. B., Wu, J. R., Pawlowsky-Glahn, V., and Egozcue, J. J. (2016), "It's All Relative: Analyzing Microbiome Data as Compositions," *Annals of Epidemiology*, 26, 322–329. [1]

Gordon-Rodriguez, E., Quinn, T. P., and Cunningham, J. P. (2022), "Learning Sparse Log-Ratios for High-Throughput Sequencing Data," *Bioinformatics*, 38, 157–163. [2]

Greenacre, M. (2019), "Variable Selection in Compositional Data Analysis Using Pairwise Logratios," *Mathematical Geosciences*, 51, 649–682. [2]

Greenacre, M., Martínez-Álvaro, M., and Blasco, A. (2021), "Compositional Data Analysis of Microbiome and Any-Omics Datasets: A Validation of the Additive Logratio Transformation," *Frontiers in Microbiology*, 12, 727398. [1]

Hong, Q., Chen, G., and Tang, Z.-Z. (2023), "PhyloMed: A Phylogeny-based Test of Mediation Effect in Microbiome," *Genome Biology*, 24, 1–21. [8]

Huang, P., Ji, F., Cheung, A. H.-K., Fu, K., Zhou, Q., Ding, X., Chen, D., Lin, Y., Wang, L., Jiao, Y., et al. (2024), "Peptostreptococcus Stomatis Promotes Colonic Tumorigenesis and Receptor Tyrosine Kinase Inhibitor Resistance by Activating ERBB2-MAPK," *Cell Host & Microbe*, 32, 1365–1379. [7]

Jobin, C. (2017), "Human Intestinal Microbiota and Colorectal Cancer: Moving Beyond Associative Studies," *Gastroenterology*, 153, 1475–1478. [7]

Lin, H., and Peddada, S. D. (2024), "Multigroup Analysis of Compositions of Microbiomes with Covariate Adjustments and Repeated Measures," *Nature Methods*, 21, 83–91. [3]

Lin, W., Shi, P., Feng, R., and Li, H. (2014), "Variable Selection in Regression with Compositional Covariates," *Biometrika*, 101, 785–797. [2,5]

McLaren, M. R., Willis, A. D., and Callahan, B. J. (2019), "Consistent and Correctable Bias in Metagenomic Sequencing Experiments," *Elife*, 8, e46923. [1]

Parker, B. J., Wearsch, P. A., Veloo, A. C., and Rodriguez-Palacios, A. (2020), "The Genus Alistipes: Gut Bacteria with Emerging Implications to Inflammation, Cancer, and Mental Health," *Frontiers in Immunology*, 11, 906. [7]

Randolph, T. W., Zhao, S., Copeland, W., Hullar, M., and Shojaie, A. (2018), "Kernel-Penalized Regression for Analysis of Microbiome Data," *The Annals of Applied Statistics*, 12, 540–566. [2]

Rubinstein, M. R., Wang, X., Liu, W., Hao, Y., Cai, G., and Han, Y. W. (2013), "Fusobacterium Nucleatum Promotes Colorectal Carcinogenesis by Modulating E-cadherin/β-catenin Signaling via its FadA Adhesin," *Cell Host & Microbe*, 14, 195–206. [3,7]

Ryu, H., Henson, M., Elk, M., Toledo-Hernandez, C., Griffith, J., Blackwood, D., Noble, R., Gourmelon, M., Glassmeyer, S., and Santo Domingo, J. W. (2013), "Development of Quantitative PCR Assays Targeting the 16S rRNA Genes of Enterococcus spp. and their Application to the Identification of Enterococcus Species in Environmental Samples," *Applied and Environmental Microbiology*, 79, 196–204. [1]

Sung, H., Ferlay, J., Siegel, R. L., Laversanne, M., Soerjomataram, I., Jemal, A., and Bray, F. (2021), "Global Cancer Statistics 2020: Globocan Estimates of Incidence and Mortality Worldwide for 36 Cancers in 185 Countries," *CA: A Cancer Journal for Clinicians*, 71, 209–249. [2]

Tang, Z.-Z., Chen, G., Alekseyenko, A. V., and Li, H. (2017), "A General Framework for Association Analysis of Microbial Communities on a Taxonomic Tree," *Bioinformatics*, 33, 1278–1285. [8]

Thomas, A. M., Manghi, P., Asnicar, F., Pasolli, E., Armanini, F., Zolfo, M., Beghini, F., Manara, S., Karcher, N., Pozzi, C., et al. (2019), "Metagenomic Analysis of Colorectal Cancer Datasets Identifies Cross-Cohort Microbial Diagnostic Signatures and a Link with Choline Degradation," *Nature Medicine*, 25, 667–678. [3]

Tibshirani, R. (1996), "Regression Shrinkage and Selection via the Lasso," *Journal of the Royal Statistical Society*, Series B, 58, 267–288. [1]

Vandeputte, D., Kathagen, G., D'hoe, K., Vieira-Silva, S., Valles-Colomer, M., Sabino, J., Wang, J., Tito, R. Y., De Commer, L., Darzi, Y., et al. (2017), "Quantitative Microbiome Profiling Links Gut Community Variation to Microbial Load," *Nature*, 551, 507–511. [1]

Vogtmann, E., Hua, X., Zeller, G., Sunagawa, S., Voigt, A. Y., Hercog, R., Goedert, J. J., Shi, J., Bork, P., and Sinha, R. (2016), "Colorectal Cancer and the Human Gut Microbiome: Reproducibility with Whole-Genome Shotgun Sequencing," *PloS One*, 11, e0155362. [3]

Wainwright, M. J. (2019), *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*, Cambridge: Cambridge University Press. [4]

Wang, T., and Zhao, H. (2017), "Structured Subcomposition Selection in Regression and its Application to Microbiome Data Analysis," *The Annals of Applied Statistics*, 11, 771–791. [2,8]

Wirbel, J., Pyl, P. T., Kartal, E., Zych, K., Kashani, A., Milanese, A., Fleck, J. S., Voigt, A. Y., Palleja, A., Ponnudurai, R., et al. (2019), "Meta-Analysis of Fecal Metagenomes Reveals Global Microbial Signatures that are Specific for Colorectal Cancer," *Nature Medicine*, 25, 679–689. [3]

Yu, J., Feng, Q., Wong, S. H., Zhang, D., yi Liang, Q., Qin, Y., Tang, L., Zhao, H., Stenvang, J., Li, Y., et al. (2017), "Metagenomic Analysis of Faecal Microbiome as a Tool Towards Targeted Non-invasive Biomarkers for Colorectal Cancer," *Gut*, 66, 70–78. [3]

Zeller, G., Tap, J., Voigt, A. Y., Sunagawa, S., Kultima, J. R., Costea, P. I., Amiot, A., Böhm, J., Brunetti, F., Habermann, N., et al. (2014), "Potential of Fecal Microbiota for Early-Stage Detection of Colorectal Cancer," *Molecular Systems Biology*, 10, 766. [3]

Zhao, L., Zhang, X., Zhou, Y., Fu, K., Lau, H. C.-H., Chun, T. W.-Y., Cheung, A. H.-K., Coker, O. O., Wei, H., Wu, W. K.-K., et al. (2022), "Parvimonas Micra Promotes Colorectal Tumorigenesis and is Associated with Prognosis of Colorectal Cancer Patients," *Oncogene*, 41, 4200–4210. [7]

Zhu, J., Wang, X., Hu, L., Huang, J., Jiang, K., Zhang, Y., Lin, S., and Zhu, J. (2022), "abess: A Fast Best-Subset Selection Library in Python and R," *The Journal of Machine Learning Research*, 23, 9206–9212. [1]